



. Load the dataset and import it into a Pandas DataFrame.

```
In [8]: from google.colab import files
import pandas as pd
import tensorflow as tf

# Upload the CSV file
uploaded = files.upload() # Choose "Real estate.csv"

# Load into Pandas DataFrame
df = pd.read_csv("Real estate - Real estate.csv") # make sure filename matches
print("◆ Dataset loaded successfully!")
print(df.head())

# Convert to TensorFlow tensor
data_tensor = tf.convert_to_tensor(df.values, dtype=tf.float32)
print("\nTensorFlow tensor shape:", data_tensor.shape)
```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving Real estate - Real estate.csv to Real estate - Real estate (1).csv

◆ Dataset loaded successfully!

	X1 transaction date	X2 house age	X3 distance to the nearest MRT station \
0	2012.917	32.0	84.87882
1	2012.917	19.5	306.59470
2	2013.583	13.3	561.98450
3	2013.500	13.3	561.98450
4	2012.833	5.0	390.56840

	X4 number of convenience stores	X5 latitude	X6 longitude \
0	10.0	24.98298	121.54024
1	9.0	24.98034	121.53951
2	5.0	24.98746	121.54391
3	5.0	24.98746	121.54391
4	5.0	24.97937	121.54245

	Y house price of unit area
0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

TensorFlow tensor shape: (415, 7)

Display the first five rows and the last three rows of the dataset

```
In [9]: print("First 5 rows:")
print(df.head())
```

```
print("\nLast 3 rows:")
print(df.tail(3))
```

First 5 rows:

	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	\
0	2012.917	32.0		84.87882
1	2012.917	19.5		306.59470
2	2013.583	13.3		561.98450
3	2013.500	13.3		561.98450
4	2012.833	5.0		390.56840

	X4 number of convenience stores	X5 latitude	X6 longitude	\
0	10.0	24.98298	121.54024	
1	9.0	24.98034	121.53951	
2	5.0	24.98746	121.54391	
3	5.0	24.98746	121.54391	
4	5.0	24.97937	121.54245	

	Y house price of unit area
0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

Last 3 rows:

	X1 transaction date	X2 house age	\
412	2013.000	8.1	
413	2013.500	6.5	
414	2013.167	1.9	

	X3 distance to the nearest MRT station	X4 number of convenience stores	\
412	104.81010	5.0	
413	90.45606	9.0	
414	355.00000	NaN	

	X5 latitude	X6 longitude	Y house price of unit area
412	24.96674	121.54067	52.5
413	24.97433	121.54310	63.9
414	24.97293	121.54026	40.5

Get the dimensions (number of rows and columns) of the dataset.

```
In [10]: print("\nShape of dataset (rows, columns):", df.shape)
```

Shape of dataset (rows, columns): (415, 7)

Generate descriptive statistics (mean, median, standard deviation, five-point summary, IQR, etc.) for the data

```
In [11]: print("\nDescriptive Statistics:")
print(df.describe(include="all"))
```

```
print("\nMedian values:")
print(df.median())

print("\nFive-point summary:")
print(df.describe().loc[["min", "25%", "50%", "75%", "max"]])

# IQR
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
print("\nInterquartile Range (IQR):")
print(IQR)
```

# Descriptive Statistics:

	X1 transaction date	X2 house age \
count	415.000000	415.000000
mean	2013.149014	17.674458
std	0.281628	11.405161
min	2012.667000	0.000000
25%	2012.917000	8.950000
50%	2013.167000	16.100000
75%	2013.417000	28.100000
max	2013.583000	43.800000

	X3 distance to the nearest MRT station \
count	415.000000
mean	1082.129338
std	1261.092057
min	23.382840
25%	289.324800
50%	492.231300
75%	1452.760000
max	6488.021000

	X4 number of convenience stores	X5 latitude	X6 longitude \
count	414.000000	415.000000	415.000000
mean	4.094203	24.969039	121.533378
std	2.945562	0.012397	0.015332
min	0.000000	24.932070	121.473530
25%	1.000000	24.963010	121.528570
50%	4.000000	24.971100	121.538630
75%	6.000000	24.977450	121.543300
max	10.000000	25.014590	121.566270

	Y house price of unit area
count	415.000000
mean	37.986265
std	13.590608
min	7.600000
25%	27.700000
50%	38.500000
75%	46.600000
max	117.500000

## Median values:

X1 transaction date	2013.16700
X2 house age	16.10000
X3 distance to the nearest MRT station	492.23130
X4 number of convenience stores	4.00000
X5 latitude	24.97110
X6 longitude	121.53863
Y house price of unit area	38.50000
dtype:	float64

## Five-point summary:

	X1 transaction date	X2 house age \
min	2012.667	0.00

25%	2012.917	8.95
50%	2013.167	16.10
75%	2013.417	28.10
max	2013.583	43.80

	X3 distance to the nearest MRT station	X4 number of convenience stores \
min	23.38284	0.0
25%	289.32480	1.0
50%	492.23130	4.0
75%	1452.76000	6.0
max	6488.02100	10.0

	X5 latitude	X6 longitude	Y house price of unit area
min	24.93207	121.47353	7.6
25%	24.96301	121.52857	27.7
50%	24.97110	121.53863	38.5
75%	24.97745	121.54330	46.6
max	25.01459	121.56627	117.5

Interquartile Range (IQR):

X1 transaction date	0.50000
X2 house age	19.15000
X3 distance to the nearest MRT station	1163.43520
X4 number of convenience stores	5.00000
X5 latitude	0.01444
X6 longitude	0.01473
Y house price of unit area	18.90000

dtype: float64

Print a concise summary of the dataset as information on data types (schema) and missing values

```
In [12]: print("\nInfo about dataset:")
print(df.info())

print("\nMissing values count:")
print(df.isnull().sum())
```

Info about dataset:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 415 entries, 0 to 414

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	X1 transaction date	415 non-null	float64
1	X2 house age	415 non-null	float64
2	X3 distance to the nearest MRT station	415 non-null	float64
3	X4 number of convenience stores	414 non-null	float64
4	X5 latitude	415 non-null	float64
5	X6 longitude	415 non-null	float64
6	Y house price of unit area	415 non-null	float64

dtypes: float64(7)

memory usage: 22.8 KB

None

Missing values count:

X1 transaction date	0
X2 house age	0
X3 distance to the nearest MRT station	0
X4 number of convenience stores	1
X5 latitude	0
X6 longitude	0
Y house price of unit area	0

dtype: int64

Add a new column named "X22" by converting the "house age" from years to days

```
In [13]: df["X22"] = df["X2 house age"] * 365  
print(df[["X2 house age", "X22"]].head())
```

	X2 house age	X22
0	32.0	11680.0
1	19.5	7117.5
2	13.3	4854.5
3	13.3	4854.5
4	5.0	1825.0

Delete the column "X22" from the dataset.

```
In [14]: df.drop("X22", axis=1, inplace=True)  
print(df.head())
```

	X1 transaction date	X2 house age	X3 distance to the nearest MRT station	\
0	2012.917	32.0		84.87882
1	2012.917	19.5		306.59470
2	2013.583	13.3		561.98450
3	2013.500	13.3		561.98450
4	2012.833	5.0		390.56840

	X4 number of convenience stores	X5 latitude	X6 longitude	\
0	10.0	24.98298	121.54024	
1	9.0	24.98034	121.53951	
2	5.0	24.98746	121.54391	
3	5.0	24.98746	121.54391	
4	5.0	24.97937	121.54245	

	Y house price of unit area
0	37.9
1	42.2
2	47.3
3	54.8
4	43.1

Create three new instances synthetically and add them to the dataset

```
In [16]: import pandas as pd

# Suppose your DataFrame is df
print("Shape before adding:", df.shape)

# Create 3 synthetic instances with the same columns
new_rows = pd.DataFrame([
    [0, 15, 560.0, 2, 24.98, 121.54, 45.0], # Example instance
    [0, 30, 1800.0, 3, 24.96, 121.52, 35.0], # Example instance
    [0, 5, 350.0, 1, 24.97, 121.50, 55.0]   # Example instance
], columns=df.columns) # use same column names

# Append to dataset
df = pd.concat([df, new_rows], ignore_index=True)

print("Shape after adding:", df.shape)
print(df.tail(5)) # show last rows including new ones
```

Shape before adding: (415, 7)

Shape after adding: (418, 7)

	X1 transaction date	X2 house age	\
413	2013.500	6.5	
414	2013.167	1.9	
415	0.000	15.0	
416	0.000	30.0	
417	0.000	5.0	

	X3 distance to the nearest MRT station	X4 number of convenience stores	\
413	90.45606	9.0	
414	355.00000	NaN	
415	560.00000	2.0	
416	1800.00000	3.0	
417	350.00000	1.0	

	X5 latitude	X6 longitude	Y house price of unit area
413	24.97433	121.54310	63.9
414	24.97293	121.54026	40.5
415	24.98000	121.54000	45.0
416	24.96000	121.52000	35.0
417	24.97000	121.50000	55.0

Delete the newly inserted three instances from the dataset.

```
In [17]: df = df[:-3]
print("After deleting new rows:", df.tail(5))
```

	X1 transaction date	X2 house age	\
410	2012.667	5.6	
411	2013.250	18.8	
412	2013.000	8.1	
413	2013.500	6.5	
414	2013.167	1.9	

	X3 distance to the nearest MRT station	X4 number of convenience stores	\
410	90.45606	9.0	
411	390.96960	7.0	
412	104.81010	5.0	
413	90.45606	9.0	
414	355.00000	NaN	

	X5 latitude	X6 longitude	Y house price of unit area
410	24.97433	121.54310	50.0
411	24.97923	121.53986	40.6
412	24.96674	121.54067	52.5
413	24.97433	121.54310	63.9
414	24.97293	121.54026	40.5

. Update the "house price of unit area" to 110, provided it is currently greater than the amount.

```
In [20]: df.loc[df["Y house price of unit area"] > 110, "Y house price of unit area"] =
```



Find the latitude and longitude of the houses whose prices are less than or equal to 20.

```
In [21]: cheap_houses = df[df["Y house price of unit area"] <= 20][["X5 latitude", "X6 longitude"]  
print(cheap_houses)
```

	X5 latitude	X6 longitude
8	24.95095	121.48458
40	24.94155	121.50381
41	24.94297	121.50342
48	24.94684	121.49578
49	24.94925	121.49542
55	24.94968	121.53009
73	24.94155	121.50381
83	24.96056	121.50831
87	24.94297	121.50342
93	24.94920	121.53076
113	24.96172	121.53812
116	24.94375	121.47883
117	24.93885	121.50383
155	24.94155	121.50381
156	24.94883	121.52954
162	24.94297	121.50342
170	24.94741	121.49628
176	24.94867	121.49507
180	24.94898	121.49621
183	24.94155	121.50381
226	24.94155	121.50381
229	24.94890	121.53095
231	24.94235	121.50357
232	24.95032	121.49587
249	24.95743	121.47516
251	24.94960	121.53018
255	24.95095	121.48458
298	24.94155	121.50381
309	24.94883	121.52954
320	24.93885	121.50383
329	24.93885	121.50383
330	24.94935	121.53046
331	24.94826	121.49587
347	24.95719	121.47353
384	24.94297	121.50342
409	24.94155	121.50381

Add the missing convenience store values of instances by calculating the average number of convenience stores

```
In [25]: # Step 1: Calculate average number of convenience stores (ignoring NaN)  
avg = df["X4 number of convenience stores"].mean()  
  
# Step 2: Fill missing values with that average  
df["X4 number of convenience stores"] = df["X4 number of convenience stores"].
```

```
# Step 3: Verify
print("Missing values after filling:")
print(df["X4 number of convenience stores"].isnull().sum())
```

Missing values after filling:  
0

Find the normalized distance to the nearest train station by performing: (a) Z-score normalization. (b) Min-max normalization. (c) Decimal scaling.

```
In [26]: x3 = df["X3 distance to the nearest MRT station"]

# (a) Z-score
z_score = (x3 - x3.mean()) / x3.std()

# (b) Min-Max
min_max = (x3 - x3.min()) / (x3.max() - x3.min())

# (c) Decimal scaling
scaling_factor = 10**len(str(int(x3.abs().max())))
decimal_scaled = x3 / scaling_factor

print("\nZ-score normalization:\n", z_score.head())
print("\nMin-Max normalization:\n", min_max.head())
print("\nDecimal scaling:\n", decimal_scaled.head())
```

Z-score normalization:

```
0    -0.790783
1    -0.614971
2    -0.412456
3    -0.412456
4    -0.548383
```

Name: X3 distance to the nearest MRT station, dtype: float64

Min-Max normalization:

```
0     0.009513
1     0.043809
2     0.083315
3     0.083315
4     0.056799
```

Name: X3 distance to the nearest MRT station, dtype: float64

Decimal scaling:

```
0     0.008488
1     0.030659
2     0.056198
3     0.056198
4     0.039057
```

Name: X3 distance to the nearest MRT station, dtype: float64

Generate the following basic visualizations using Seaborn. Customize your visualizations by adding titles, labels, legends, and appropriate color schemes. (a) Create a histogram for the "Y house price of unit area" attribute.


(b) Create a box-and-whisker plot for the "Y house price of unit area" attribute. (c) Create a scatter plot showing house prices against house age. (d) Add a second scatter plot showing house prices against distance to the nearest MRT station.

```
In [27]: sns.histplot(df["Y house price of unit area"], kde=True, color="skyblue")
plt.title("Histogram of House Price per Unit Area")
plt.show()


sns.boxplot(y=df["Y house price of unit area"], color="lightcoral")
plt.title("Boxplot of House Price per Unit Area")
plt.show()

sns.scatterplot(x=df["X2 house age"], y=df["Y house price of unit area"], color="lightcoral")
plt.title("House Price vs House Age")
plt.show()


sns.scatterplot(x=df["X3 distance to the nearest MRT station"], y=df["Y house price of unit area"], color="lightcoral")
plt.title("House Price vs Distance to MRT station")
plt.show()
```




Boxplot of House Price per Unit Area



House Price vs House Age





Form the Design Matrix  $X$  of shape  $m \times n + 1$  in order to apply normal equation method where  $m$  is the number of training examples and  $n$  is the number of input features. Only use the two normalized input features 'X2 house age' and 'X3 distance to the nearest MRT station' from the dataset as second and third columns respectively and all 1 s as the first column. Also, form output vector  $Y$  of shape  $m \times 1$

```
In [28]: # Using normalized features (Min-Max)
x2 = (df["X2 house age"] - df["X2 house age"].min()) / (df["X2 house age"].max() - df["X2 house age"].min())
x3 = (df["X3 distance to the nearest MRT station"] - df["X3 distance to the nearest MRT station"].min()) / (df["X3 distance to the nearest MRT station"].max() - df["X3 distance to the nearest MRT station"].min())

m = len(df)
X = np.c_[np.ones(m), x2, x3] # m x (n+1)
Y = df["Y house price of unit area"].values.reshape(-1, 1) # m x 1

print("X shape:", X.shape)
print("Y shape:", Y.shape)
```

X shape: (415, 3)

Y shape: (415, 1)

Find the parameter vector  $W$  using the normal equation method as  $W = (X^T X)^{-1} X^T Y$ .

```
In [29]: W = np.linalg.inv(X.T @ X) @ X.T @ Y
print("Parameters (W) from Normal Equation:\n", W)
```

```
Parameters (W) from Normal Equation:
[[ 49.61767979]
 [ -9.99718231]
 [-46.49889746]]
```

Implement the gradient descent algorithm with the following steps.

- Form the Design Matrix  $X$  of shape  $n \times m$ . Only use the two normalized input features 'X2 house age' and 'X3 distance to the nearest MRT station' and the output vector  $Y$  of shape  $1 \times m$ .
- Initialize the parameter vector  $W$  of shape  $1 \times n$  and bias  $b$  (scalar).
- Repeat the following steps to a certain number of iterations with learning rate  $\alpha = 0.01$ , and print the final parameter values.
  - Calculate the prediction  $\hat{Y} = W X + b$ .
  - Compute loss  $L = \frac{1}{2} \times (Y^{\wedge} - Y)^2$
  - Compute error  $E = Y^{\wedge} - Y$
  - Compute the gradient with respect to  $W$  as  $dW = \frac{1}{m} E \cdot X^T$  and with respect to  $b$  as  $db = \frac{1}{m} \times E$  (sum over the columns)
  - Update  $W = W - \alpha dW$  and  $b = b - \alpha db$
- Use tensorflow GradientTape() to automatically calculate the gradients in the above step (d) and redo the training steps and print the final parameter values.

```
In [30]: alpha = 0.01
iterations = 1000

# Initialize
W = np.zeros((1, 2)) # weights for x2, x3
b = 0.0

X_gd = np.vstack([x2, x3]) # shape (2, m)
Y_gd = Y.reshape(1, -1) # shape (1, m)
m = Y_gd.shape[1]

for i in range(iterations):
    Y_hat = np.dot(W, X_gd) + b
    E = Y_hat - Y_gd
    dW = (1/m) * np.dot(E, X_gd.T)
    db = (1/m) * np.sum(E)
    W -= alpha * dW
    b -= alpha * db

print("Final parameters (manual GD):")
print("W:", W, " b:", b)

# TensorFlow GradientTape
W_tf = tf.Variable([[0.0, 0.0]])
b_tf = tf.Variable(0.0)

X_tf = tf.constant(X_gd.T, dtype=tf.float32) # (m, 2)
Y_tf = tf.constant(Y, dtype=tf.float32) # (m, 1)

optimizer = tf.keras.optimizers.SGD(learning_rate=0.01)
```

```

for epoch in range(1000):
    with tf.GradientTape() as tape:
        Y_pred = tf.matmul(X_tf, tf.transpose(W_tf)) + b_tf
        loss = tf.reduce_mean(tf.square(Y_pred - Y_tf))
        grads = tape.gradient(loss, [W_tf, b_tf])
        optimizer.apply_gradients(zip(grads, [W_tf, b_tf]))

print("Final parameters (TF GD):")
print("W:", W_tf.numpy(), " b:", b_tf.numpy())

```

```

Final parameters (manual GD):
W: [[ 3.33859658 -10.38361448]] b: 37.78556232358561
Final parameters (TF GD):
W: [[ -2.3015294 -21.340294 ]] b: 42.05744

```

Define a class to create a Linear Regression model with methods fit and predict. Use the above iterative process to implement the model's training within the fit method.

```

In [31]: class MyLinearRegression:
    def __init__(self, lr=0.01, epochs=1000):
        self.lr = lr
        self.epochs = epochs

    def fit(self, X, Y):
        m, n = X.shape
        self.W = np.zeros((1, n))
        self.b = 0.0
        for i in range(self.epochs):
            Y_hat = np.dot(self.W, X.T) + self.b
            E = Y_hat - Y.T
            dW = (1/m) * np.dot(E, X)
            db = (1/m) * np.sum(E)
            self.W -= self.lr * dW
            self.b -= self.lr * db

    def predict(self, X):
        return np.dot(self.W, X.T) + self.b

# Example usage
model = MyLinearRegression(lr=0.01, epochs=1000)
model.fit(X[:,1:], Y) # use only features (skip bias column)
preds = model.predict(X[:,1:])
print("Predictions shape:", preds.shape)

```

```

Predictions shape: (1, 415)

```